



Establishing AgentOps for enterprise AI on Red Hat AI

Younes Ben Brahim

Principal Product Marketing Manager
AI Business Unit

Roberto Carratalá

Principal AI Architect –
AI Business Unit


Peter Double

Principal Product Manager
AI Business Unit


Agenda

- ▶ **The agentic AI era:** Why agents are the next frontier
- ▶ **The production gap:** Why MLOps isn't enough – introducing AgentOps
- ▶ **Red Hat AI:** The safest place for enterprise agents
- ▶ The pillars powering agentic AI on Red Hat AI
- ▶ **Demo:** Agentic AI in action
- ▶ The **Red Hat AI advantage** and key takeaways

From ChatBot to Agents



The Dawn of AI Agents



The evolution of enterprise AI

How it started

Where it is!

Where it is going...

LLM **Workflow**

Agentic **Workflow**


Autonomous **Agents**



The LLM is a simple worker,
and autonomy is **Low**.
Orchestration is external.



The LLM has decision
making power, and autonomy
is **Medium**.



Given a high-level objective,
the agent defines the tasks
and plans. Autonomy is **High**.

Agentic AI is no longer experimental

it's operational, scaling fast, and redefining how work gets done.

80%

**Fortune 500
adoption**

of Fortune 500
companies already
use active AI agents.

– Microsoft Cyber
Pulse, 2026

\$632B

**AI spending
by 2028**

Worldwide spending
on AI solutions will
surpass \$632 billion
by 2028, growing at
29% CAGR.

– IDC, 2025


82%

Executive intent

Tech executives plan
to integrate AI
agents within 1–3
years, but <10% have
production-grade
deployments.

– Capgemini,
2025

Production gap



Three critical gaps: Pilot to production

Why 82% plan to deploy agents – but fewer than 10% have

Agent identity


Agents need secure, cryptographic identities to access sensitive tools under least-privilege principles – hardcoded keys are unacceptable in production

Ungoverned autonomy

Unconstrained dynamic code execution creates security risks, requiring isolated environments with deep execution tracing to monitor autonomous actions

Scalability and performance

Multi-agent architectures generate unpredictable, concurrent inference spikes that traditional infrastructure cannot handle



Red Hat AI - Closing the Gaps


Getting you to production



Agent Identity & Security

Every agent gets its own identity


- ▶ **Cryptographic Identity (SPIFFE/SPIRE):** Short-lived, keyless
- ▶ **MCP Gateway:** per-tool authorization & Audit
- ▶ **Scoped OAuth2 token exchange:** least-privilege access



Safety & Observability

Isolate, trace, evaluate every action

- **Agent Sandbox:** zero-trust container, microVM isolation.
- **MLFlow Tracing:** every LLM call, tool use, and decision logged
- **EvalHub Real-time Evaluation:** continuous scoring, drift detection



Scalability & Performance

Infrastructure that handles agentic workloads


- **vLLM:** high-throughput model serving
- **Llm-d:** dynamic inference routing
- **Inference-aware autoscaling:** scales on KV cache pressure

AgentOps with Red Hat AI

Keeping you in production


What is it:

A unified approach for managing the full lifecycle, security, and performance of AI agents and their tools




Why it matters:

Provides operational rigor to reliably move agents from experimentation to production with **platform-native lifecycle management**.



The complete stack metal to agents

Streamlined Agent Operationalization



AgentOps

Demo

red.ht/agentops-demo




The Red Hat AI Advantage






Focus areas

Accelerate the development and delivery of AI solutions across hybrid-cloud environments



Governance and security from metal to agents

Run any model or agent across any hardware in any cloud environment



Red Hat AI




Inference

Increase efficiency with **fast, flexible and efficient** inferencing



Data

Simplified experience for **connecting data to models and agents**



Agents

Scale Agentic AI with governance and control

Platform


Governance and security from metal to agents

Run any model or agent across any hardware in any cloud environment



Red Hat AI Enterprise

An integrated AI platform for enterprise AI deployments



AI services layer: build, tune, and manage all AI models and agent workflows.

Platform services layer: core foundation for consistency, security, and scalable operations.

Any hardware: support for various hardware accelerators for optimized performance and cost.

Deploy anywhere: choose where to train, tune, deploy and run AI models, agents and applications.

The products that power **Red Hat AI**

Red Hat AI meets customers where they are in their adoption journey providing choice of models, hardware, and location—whether on premise, in the cloud, or at the edge.



Red Hat AI Enterprise

For enterprise AI, deploy and scale efficiently, anywhere



Red Hat AI Inference

For optimized inference of large language models

Key Takeaways

What to remember from
this session

Learn more, and try out, Red Hat AI

Accelerate your AI journey
with the Red Hat
AI quickstarts


red.ht/blog-ai-quickstarts




AI quickstarts home page:
red.ht/ai-quickstarts

Subscribe to the
Red Hat AI blog


red.ht/rhai-blog



Follow us on X and YouTube



x.com/RedHat_AI



youtube.com/@RedHat

Interactive demos

- red.ht/InferenceBenchmark
- red.ht/intro-inference
- red.ht/model-optimization
- red.ht/model-customization
- red.ht/intro-llm-d
- red.ht/MaaS

60-day trial of Red Hat AI Inference

red.ht/ai-inference-server-trial



60-day trial of Red Hat Enterprise Linux AI

red.ht/RHELAI-trial



60-day trial of Red Hat OpenShift AI

red.ht/RHOAI-trial



The Red Hat AI team wants to hear from *you!*



Let us know your thoughts on Red Hat AI by
signing up to participate in research opportunities!

Scan the QR code or go to:
red.ht/uxfeedback

Have additional questions about AgenticOps?
[Visit us at the Red Hat AI Booth](#)

AI quickstarts

Explore a variety of innovative AI use cases from our community. These simple, focused examples have been designed for fast and easy deployment on the Red Hat AI platform.

 This content is authored by Red Hat experts, but has not yet been tested on every supported configuration.



red.ht/AI_quickstarts

Audit AI apps to meet compliance goals

Deploy a financial audit system with insights into the performance, cost, and usage of your large language models in real time.



Build an AI-powered virtual agent

Build and deploy a conversational AI virtual agent on Red Hat OpenShift AI to automate customer interactions and provide instant support.



Centralize company knowledge with an Enterprise RAG Chatbot

Use retrieval-augmented generation (RAG) to enhance large language models with specialized data sources for more accurate and context-aware responses.





Red Hat AI



Thank you



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



[@redhatsummit](https://twitter.com/redhatsummit)



[@redhat](https://www.instagram.com/redhat)

Red Hat AI - Closing the Gaps

Getting you to production

Agent Identity & Security

Every agent gets its own identity

- ▶ **Cryptographic Identity (SPIFFE/SPIRE):** Short-lived, keyless
- ▶ **MCP Gateway:** per-tool authorization & Audit
- ▶ **Scoped OAuth2 token exchange:** least-privilege access



Ungoverned Autonomy

Isolate, trace, evaluate every action

- **Agent Sandbox:** zero-trust container, microVM isolation.
- **MLFlow Tracing:** every LLM call, tool use, and decision logged
- **EvalHub Real-time Evaluation:** continuous scoring, drift detection



Scalability & Performance

Infrastructure that handles agentic workloads

- **vLLM:** high-throughput model serving
- **Llm-d:** dynamic inference routing
- **Inference-aware autoscaling:** scales on KV cache pressure



AI quickstarts

Explore a variety of innovative AI use cases from our community. These simple, focused examples have been designed for fast and easy deployment on the Red Hat AI platform.

 This content is authored by Red Hat experts, but has not yet been tested on every supported configuration.



red.ht/AI_quickstarts

Audit AI apps to meet compliance goals

Deploy a financial audit system with insights into the performance, cost, and usage of your large language models in real time.

Build an AI-powered virtual agent

Build and deploy a conversational AI virtual agent on Red Hat OpenShift AI to automate customer interactions and provide instant support.

Centralize company knowledge with an Enterprise RAG Chatbot

Use retrieval-augmented generation (RAG) to enhance large language models with specialized data sources for more accurate and context-aware responses.