

Note

Tue, May 12, 2026 2:32PM 11:35

SUMMARY KEYWORDS

AI automation, ops manager, VMs management, Ansible automation, root cause analysis, AI agent, developer experience, IT ops experience, ticketing system, model as a service, context engineering, agent skills, Kira, Rocketchat, DevOps.



SPEAKERS

Speaker 5, Speaker 4, Speaker 1, Speaker 3, Speaker 6, Speaker 8, Speaker 7, Speaker 2

S Speaker 1 00:20
Get the party started.

S Speaker 2 00:27
Okay, thank you. Thank you for choosing this session. We're going to kick it off. You'll be very glad to know that we really only have one slide for you. Everything's in the lab, but it will be interactive. And you know, we can pull stuff up our pathway through if there's lots of questions. So my name is Tony Kay. I'm an AI leader at for a group called portfolio technology, which probably doesn't mean a lot to you, but we're the people who are actually running this event. So we actually run the infrastructure and we cover the whole portfolio.

S Speaker 3 01:00
So that's what I do. My

S

Speaker 4 01:02

name is Bert Sutter, and I focus primarily at Red Hat on developer experience. And this is not that sort of layout here. We're going to be focusing on IT ops experience, which I'm actually very, very excited to be part of, because you guys teach me a ton of things based on having that perspective. Right? What does it mean to be on the upside house receiving those hand grenades from the development team and have you basically automate things that might make your life easier and better, because that's what this is all about.

S

Speaker 5 01:31

My name is Ragu. I'm part of North America tech sales. Ai specialist. I say I focus on FSI customers, banks, insurance. I'm glad to support you all and enjoy the lab.

S


Speaker 2 01:47


And we have a mystery guest. We have brought along a real ops manager, because this is an ops centric AI, magenta cloud, but it's about ops. So we brought along an ops manager. We're lucky to have him, because right now he's responsible for somewhere in the region of 8000 to 10,000 VMs, which are running this event, everything you do, everything you're doing now, and that's not including the AWS stuff. You can probably tell that he owns the entire O'Reilly Carl publications of the 20th century. We don't hold that against it. Called free lab.

S

Speaker 4 02:28

Now I'm not gonna play the game of you know, how many VMs you manage? Anything, ladies, anyone got more than 8000 VMs to manage. We have some customers like that. Okay, fantastic, right? We got some big numbers out there. How many people more in that 100 to 500 range between VMs and containers to manage? Okay, that's the kind of the majority group. And then how many people really beyond you know that 5000 let's say 500 to 2000 range. How many people in that range? Okay, all right. That gives me a feel for what folks are doing. And the reason I asked that basic question is, often, come to Summit, and I will ask that question of people that I meet at Summit, and I have to say, 10 years ago, it often the number was often I've got 200 for my entire business, or I've got 300 for my entire enterprise. And now we're starting to see numbers that are often greater than 1000 and that's one operations with a fleet of 1000 things to operate. Let's say, hopefully that kind of sinks out with what you guys are thinking. Also, is that fair coming from you?





S Speaker 6 03:31

How many of you build and delete 6000 plus VMs in a day? Okay, you and I can probably talk, because that's the kind of stuff that we have to do to make this event work. How many of you are less than 100 VM changes in a day? How about less than 1000 try to figure the ranges. We're okay. So you got people more than that? Bert, wow. I

S Speaker 4 03:56

want to ask you guys a couple other questions, if you guys want to indulge me, because, again, I come from the Development The development side of the house, and I'm trying to understand the world view that you guys have. There's going to be a lot of people, I imagine, that are trying to sell you an AI agent at this point. As matter of fact, we're not trying to sell you one. We're giving you one that's part of the workshop. It's free and it's open source, and you're going to get a chance to actually spend some time with it. But how many people have already looked at the pitch, if you will, from your service now, from your Atlassian, from your data dog, etc, etc, have been pitched an AI agent that help you automate things. How many people seen that sort of thing already that at least that sales pitch? Okay, how many people have actually purchased one of these automatic, automatic agent, things to help them. Wow, there's not that many. Okay, I think that's interesting by itself,

S Speaker 6 04:48

from an ops perspective. It scares the crap out of me. He

S Speaker 4 04:54

already gave us that feedback about our demo agent. Also, Tony, what do you think about

S Speaker 2 05:01

that, about what

S

Speaker 4 05:02

the use of agents in ops?

S

Speaker 2 05:04

So I think that all of you want to quadruple the number of tickets. You have to close and you're going to be thanking us, because we're going to generate you so many tickets, but they are going to be root

S

Speaker 4 05:17

caused.

S

Speaker 2 05:18

I think that's good news, because then you can quadruple your productivity.

S

Speaker 4 05:22

Yeah. So if you look at the diagram here, just so you guys understand the nature of the lab, it's going to start with Ansible automation, controller automation platform, right? So in other words, a job in there that fails becomes a trigger, the triggering event. It's look for failed jobs within Ansible, it's going to run Athena, in this case, which is the custom agent that's created. You're going to have a chance to mess with and modify yourself. And it's just a little open source thing. It's an agent built with Markdown. Might be a way to think of it, right? So it's not something you have to write a lot of code for. It's an agent built with Markdown, which is a new way of building agents in 2026 and then from there, it's going to basically do the root cause analysis, trying to figure out what to do with this thing. And has a team of sub agents. Sub agents, which you guys got a chance to see, and then it's route that to the ticketing system update you on rocket chat. So what we've done is we didn't have slack for everybody, but we have rocket chat to simulate that chat ops experience, having to interact with something. Is that cool? Is that fair? How many people here on Slack already, though? Okay, not everybody. What other chat systems are you guys using for your

S

Speaker 6 06:27

zoom? What's that? Zoom,

 Speaker 4 06:28
zoom,

 Speaker 1 06:29
okay,

 Speaker 5 06:30
yeah, so

 Speaker 4 06:32
it seems like there's quite a variety in this case. All right, our ticketing system is not JIRA. You guys notice the name change there Kira? Well, that would have been too hard to integrate everyone into our JIRA. We run massive JIRA at Red Hat, but in this case, how about a five coded Kira? It's got a cooler name. It's batter, it's batter, and Tony is the creator of Kira. So you guys are gonna get a chance to work with that also, there's also a model as a service that you're gonna get chance to interact to interact with, because we are going to be using AI models extensively throughout the little workshop you were working on, including the changing of models on the fly, just so you can see different models and how their behavior changes the application, in this case, the agent, and how it's troubleshooting tickets. But the models actually run on a model as a service that basically is running on top of ownership. Ai,

S

Speaker 2 07:23

okay, there's one other thing I did say, only one slide, but also one diagram, just so you know what happened. So Byron introduced that and just so so. AP, two controller is just a metaphor. It could be it could be Argo CD it could be multiple Argo CDs. It could be whatever runs your TerraForm, their sources, sources of errors. And what happens is, when the deep agents team gets deployed, you don't have to change your infrastructure. What it does is it hooks in, in the case of automation controller, how many of you run for lists, okay, so it just hooks in and says, whenever you have an error, tell me about it. You don't have to give me the error. You just give me the job number and tell me. And then the web hook fires the deep agent. There's only one deep agent, the ops manager. It's called catches that classifies it, and once he classifies it. When he catches it, he retrieves it, all the artifacts, error logs, what we call projects, what we call templates, etc. He pulls all that in. He has a very tight, beautiful context. And this is something important. We don't have a lot of time to study the context, but if you give a really good context to a great open source model, you will get a good result. If you give an awful context, you will not so context engineering is a thing, and so classifies it and then creates dynamically an SRE stop aging with skills. How many people have heard of agent skills? Anthropic skills? Okay, right? So these agents, that's where they get their ability. They have skills, and their skills include two things, how do I debug SSH? How do I debug DNF? But they also include institutional or tribal knowledge. If you do something strange in your environment like we do, you can capture that in the skill pertaining to that thing, and that will help the LLM. It will give it a huge clue, because it knows, not only are you using SSH, but you do it slightly differently from everybody else. Then they generate a root cause analysis. They create a ticket in Kira and they inform your Slack channel, or rocket chat, in our case. So that is all you really need to know. Everything else is in the lab. We're probably going to interrupt you at various times to get feedback and whether this is good, whether this is bad, whether it's valid, whether it relates to your world. And here is the lab, basically. Some of them are calm on your laptops. We'll leave this slide up. This slide is going to be up for quite a bit. We might throw some other slides up. You're going to select yours. It might not be at that particular point in the grid, lb, 2645, agentic, DevOps, and your password is DevOps, three. Okay, so you would you, and then that would take you straight into the mouth. I'm

S

Speaker 6 10:39

sorry. Would you say for email address?

S

Speaker 4 10:41

Yeah, your your normal

S

Speaker 6 10:42

mind?



Speaker 4 10:43

Yeah. It's just your normal email address. That's the unique identifier, so we know who you are, logged in and a DevOps three, normal three should get You right in anyone's



Speaker 7 11:00

having issues, let us know. Publish,



Speaker 8 11:13

option, move into what.