


Note

Wed, May 13, 2026 9:31AM 7:16

SUMMARY KEYWORDS

Agentic AI, Agent ops, AI business unit, Autonomous agents, Inference challenges, Data focus, Scalability, Agent identity, Governance, Zero trust, Egress control, Binary isolation, Observability, MLFlow, Agent lifecycle.



SPEAKERS

Speaker 3, Speaker 1, Speaker 2

S

Speaker 1 00:00

Agentic AI and Agent ops, and what this lab will be about, and we'll dive right into it. My name is Taylor Smith. I am a Technical Marketing Manager in our AI business unit. And this is Rebecca.

S

Speaker 2 00:12

Hello. I'm Roberto gradula. I'm also working in the AI business unit as an AI actor. I

S

Speaker 1 00:18

don't know I keep walking. Just gonna stand on this side that's safer. Okay. So, oh, and Tony's here as well. Tony say, wave, hi. He's our

S


Speaker 3 00:29

helper today, and

S

Speaker 1 00:31

he helped make sure all these instances came up healthy for all of you to work with. So we are clearly in kind of the energetic, you know, hype air era bubble, where we're moving from traditional, large language model deployments and workflows to people starting to experiment and deploy agents. And it's happening, you know, slowly and carefully, hopefully in most in most cases, especially with enterprise customers. And we're hopefully taking kind of baby steps there where we're starting to maybe give an agentic application access to certain tools and certain external systems, and maybe there's some sense of autonomous behavior. But we're not necessarily all the way into that autonomous agents era yet, but certainly we're talking about it and we're approaching it, and that introduces a lot of new challenges and considerations that we have to take into account. But this is definitely where we're going, and the lab today will go over tools that we have in place today in our products that help you kind of secure and observe these agentic applications in production. This is not no longer just, you know, an idea or theoretical. This is, you know, getting adopted at a large scale in many, many companies. Again, it's maybe smaller use cases that we're starting out with. But it's definitely no longer just an idea in our head, it's happening now, and we need to think critically about it today. So in terms of our platform, this is kind of our pillar focus slide that shows what our focus areas are. In the Red Hat AI platform inference is a huge focus area for us, and extremely important when we think about agentic workflows and agentic applications. Agents are taking a lot of different turns and leveraging a lot more tokens than a traditional large language model deployment and a traditional AI workload, which means efficient inference becomes increasingly challenging and increasingly important. Data is an important focus area as well, because we want to make sure that our agents have the necessary business context from our data and aren't just using, you know, their training data without that necessary enterprise context. And then making sure that we can deploy agents at scale, securely and be able to observe those agents. These are all extremely important areas for us in terms of product focus, and then the underlying platform that takes care of your AI workloads, from the metal layer all the way up to the agentic workload layer. Now when we start to deploy these agents in production, like I said, This introduces a lot more complexity and challenge that we didn't necessarily have to think of with traditional AI workloads, and that a lot of people today I'm seeing are not maybe don't even know what to think about, or don't know how to approach handling these new concerns. So one of those is Agent identity, so making sure that you understand who your agent is, what your agent can access and what it can't access, and assigning the necessary identity to that agent, the challenge of ungoverned autonomy of the agent. So we're trying to have these autonomous agents that can access all these different types of systems that requires a whole new layer of governance to ensure that that is happening safely and securely at scale, and then the scalability and performance layer, like I mentioned, agents have a lot of different turns. They're doing a lot of different prompt calls just in a single workflow, which increases the inference load dramatically. So you need to make sure that you have the necessary infrastructure at the foundation to handle that at scale.



The LLM is a simple worker, and autonomy is **Low**. Orchestration is external.

The LLM has decision making power, and autonomy is **Medium**.

Given a high-level objective, the agent defines the tasks and plans. Autonomy is **High**.



The agentic AI explosion

Agentic AI is no longer experimental

it's operational, scaling fast, and redefining how work gets done.

- 80% Fortune 500 adoption**
of Fortune 500 companies already use active AI agents.
— Microsoft Cyber Pulse, 2026
- \$632B AI spending by 2028**
Worldwide spending on AI solutions will surpass \$632 billion by 2028, growing at 29% CAGR.
— IDC, 2025
- 82% Executive intent**
Tech executives plan to integrate AI agents within 1–3 years, but <10% have production-grade deployments.
— Capgemini, 2025



Red Hat AI

**Inference**
Increase efficiency with fast, flexible and efficient inferencing

**Data**
Simplified experience for connecting data to models and agents

**Agents**
Scale Agentic AI with governance and control

Platform
Coverage and security from metal to agents
Run any model or agent across any hardware in any cloud environment




Agentic AI

Three critical gaps: Pilot to production

Agent identity
Agents need secure, cryptographic identities to access sensitive tools under least privilege principles. Hardcoded keys are unacceptable in production.

Un-governed autonomy
Uncontained dynamic code execution creates security risks, requiring isolated environments with deep execution tracing to monitor autonomous actions.


Scalability and performance
Multi-agent architectures generate spindly tails. Concurrent inference scales, but hardware infrastructure cannot handle



S

Speaker 2 04:27

So we are using this streamline agent operations that we already talked about it, but we have security, safety and observability. So inside of the agent ops, we want first to govern what agent can access. So we want to give every single agent their own cryptographic identity. And for that, we are using one standard that is spiff and spirit, that is giving the short lived different credentials, and also all things like MCB gateway and the scope access, then we want to isolate and contain the agent execution. So we are using this zero trust unboxing into, for example, the different agents. And also we are adding things like egress control to not do certain things. And also we want to do this binary isolation, so we are adding open shell for filtering the network and system. So all of this is the agent sandbox and we wanted to control. And this also was natively included in OpenShift, so we are relying on certain things, but also including our others like open shell. And last but not least, you will see on this lab, and we will focus on the observability. So you want to have this monitor, trace and manage the lifecycle, having the mlflow tracing the agent lifecycle as well, and lab scoring. So we will focus on all of the end to end observability. And we are using heavily mlflow Because it's great and it's very, very flexible when you want to bring your own agent. So essentially, as Taylor explain it, we are using this that is, observe, trace and evaluate, enforce Identity and Access Control, and then deploy the different agents and manage the life cycle of this. So we have this mix, the gap between the experimental AI, two agents in production. So we have these developments to production, and we do tracing evaluations, identity section, security, observability and much more. So again, these three pieces that is observed, phrase and evaluate. You will focus on this lab, heavy on this, and also enforce Identity and Access Control with all of the things that we already explained. And then import your agent and manage the different life cycle of the different agent. And I think that with that, we can move forward. Hope that you enjoy that.





Speaker 1 07:11

So this is your lab access information, so you'll find this agent ops and production time.