Red Hat Summit

AnsibleFest

The AI trust barrier: Why and how enterprises must govern autonomous solutions

Anshul Behl
Principal Technical Marketing Manager
Red Hat Ansible

Ruchika Pandey
Principal Software Engineer
Cisco

Cisco Red Hat

What you'll walk away with

1. How enterprises can adopt AI with confidence by closing the trust gap
2. The governing principle: separate AI intelligence from AI execution
3. How Cisco AI Defense and Ansible Automation Platform work together to enable governed AI

AI is accelerating faster than governance can keep up

The capability is here

- Chatbots (2023). RAG (2024). Agentic AI (2025+). Each generation grants AI more autonomy over enterprise systems.
- 1.9M+ models on HuggingFace. 450K+ datasets. A growing ecosystem of MCP servers. The AI supply chain is largely unvetted.
- Every layer of an AI application introduces new risk: prompt injection, data poisoning, agent goal hijacking, supply chain compromise, and compliance failures.

But governance is not keeping pace

CISOs and compliance teams are asking questions that demand answers:

- *"Who approved that action?"*
- *"Can we audit this decision?"*
- *"What happens when the AI is wrong?"*

Consequences of unmanaged AI risk: financial damages, litigation, reputational harm, noncompliance, IP leakage.

EU AI Act 2024 mandates external audits of generative AI systems throughout their lifecycle.

83% plan to deploy agentic AI **29%** feel prepared to do so securely *State of AI Security 2026, Cisco*

AI decides what to do.
Automation governs how it gets done.
Trust requires securing both layers independently.

Separate AI intelligence from AI execution

Cisco AI Defense: Coverage across the AI lifecycle

Discovery

AI Cloud Visibility

Identify AI assets

Inventory the AI models, agents, and connected data sources across distributed environments to understand usage and gauge risk.

Detection

AI Supply Chain Risk Management

Scan for threats

Scan model files, repos, and MCP servers to proactively block malicious or unsafe AI assets before operations are impacted.

AI Model & App Validation

Detect the vulnerabilities

Identify safety and security vulnerabilities across models at scale with algorithmic red teaming technology.

Protection

AI Runtime Protection

Mitigate threats in real time

Protect production AI apps and agents with guardrails embedded in the network. Block attacks and harmful responses in real time.

Why this matters: AI applications are complex and non-deterministic. Each layer (User, Application, Model/Agent, Data, Infrastructure) introduces new risk vectors: business and reputational harm, data security and privacy breaches, supply chain vulnerabilities, cyber attacks, and compliance failures.

Unified management with Cisco Security Cloud Control. Backed by Cisco AI Security and Safety Framework: 20+ objectives, 150+ techniques, mapped to OWASP, MITRE ATLAS, and NIST.

CISCO Red Hat

DEMO
Cisco AI Defense

- LLM Security: https://leaderboard.aidefense.cisco.com
- Algorithmic Red Teaming: https://explorer.aidefense.cisco.com
- Open Source Tools: https://github.com/cisco-ai-defense

Red Hat

Your automation library

One security boundary, one audit trail, one source of truth

Your automation library is also your security boundary.

It should define what AI is - *and isn't* - allowed to do.

In the AI era, agents will need a "menu" of approved actions.

Your automation library IS that menu!

The trusted execution layer for AI-driven IT operations

Task-based

Operational Knowledge

Scripts – Playbooks – Runbooks

Event-driven

Operational Data

Observability · Monitoring · Analytics · Event streams

AI-driven

AI / Inference

Predictive analytics · Agents · Natural Language Processing

Red Hat Ansible Automation Platform

RBAC · Approval Gates · Auditing · Content Signing · Credential Management

Unified governance across IT operations

IT infrastructure

Red Hat

The governance infrastructure AI needs is already running in production

Red Hat Ansible Automation Platform delivers the complete stack

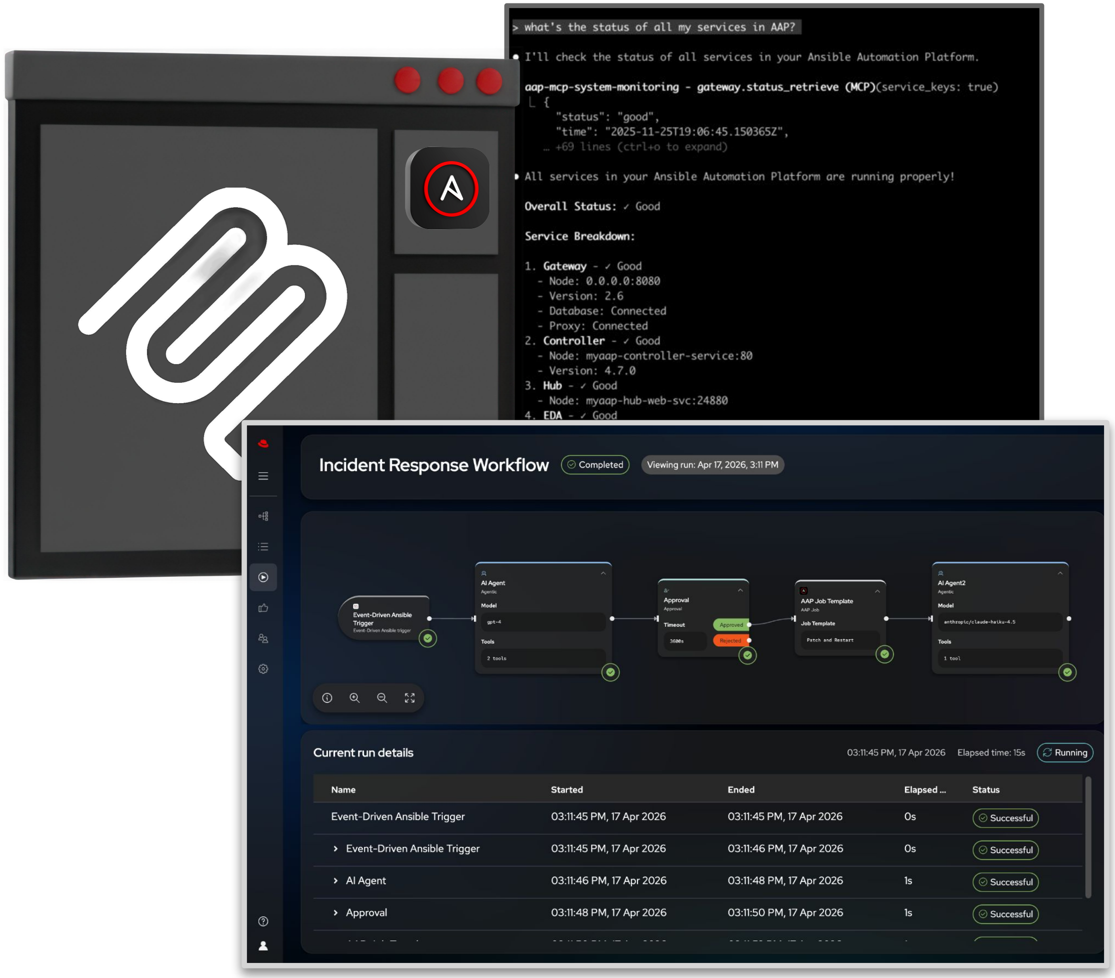

- ✓ AAP plugin framework for MCP to enable **task-driven automation** for MCP actions
- ✓ Red Hat Ansible Automation Platform MCP Servers that enable **event-driven AI Ops automation**
- ✓ New automation workflow tooling that enables **AI-driven automation** to be orchestrated in concert with task and event-driven automation.
  - Side-by-side execution of multiple automation modes
  - Context flows across modalities
- ✓ Current AAP agent-driven investments establish a framework where the future combines the guardrails and compliance standards of AAP with fully autonomous agents.

DEMO
AI-Driven Incident Response

Agents + MCP + Ansible Automation Platform

Red Hat

Key takeaways

Intelligence without execution is just a recommendation.

Cisco AI Defense secures the intelligence layer. Ansible Automation Platform turns AI decisions into governed, auditable actions.

Your automation library is your governance boundary.

AI agents interact through MCP with only the playbooks they are authorized to use. Every action is deterministic, repeatable, and fully auditable.

Separate intelligence from execution to build trust.

Cisco AI Defense validates and protects AI applications. Ansible governs what they are allowed to do. Together: from cautious skepticism to confident, governed deployment.

Red Hat Summit

Get started today

- AIOps overview: redhat.com/automation-aiops
- Market research paper: red.ht/futurum_redhat_aiops
- The State of AI security report(Cisco)
- Cisco AI Defense repo- Apache 2.0 (Clone it today)

Cisco Red Hat