


Note

📅 Thu, May 14, 2026 8:46AM ⌚ 44:17

SUMMARY KEYWORDS

AutoML, AutoRag, AI expertise, feature engineering, hyperparameters, data science, predictive AI, generative AI, synthetic data, evaluation data, MLflow, automation, enterprise AI, knowledge discovery, document retrieval.



SPEAKERS

Speaker 8, Speaker 3, Speaker 4, Speaker 6, Speaker 5, Speaker 7, Speaker 10, Speaker 9, Speaker 2, Speaker 1


S




Speaker 1 00:00

Particular we're going to be talking about these outlay the patterns with three cores, which, by the way, was released today. You will have the first two. One that is outgo. Now that is not the automl library. Is automl as the pattern and outer right? But this is something that in red AI, we will continue to bring the outdoor experiences all throughout the product. So today we're gonna focus on this too. And so one of the things that we have seen is that in the enterprise, these two are extremely important, right? Traditional lives are going away, and is there, but the challenges are really my chemicals and generated AI part is okay, so copilot agents. We don't have enough experts. We have a lot of people that like tinkering a little bit here and there, following tutorials. But what is the right combination after I do the POC, and how to find that? And that's when we really need some expertise. And for example, in traditional AI, if you come from traditional AI, you will know about feature engineering, which is everything that I need to extract from my data, put in a fancy structure data set, and determine from there which one are important to predict, whatever I want to predict. And that is, it's a tedious process at the beginning and then maintaining that every time there are drift over time is you need expert, not only SME, expert on what that thing is trying to help with, but also in the data scientist side. And then it comes not only selecting the right model architecture for the traditional learning, but also going and setting up and finding out all those hyper parameters when you actually look at the test matrix of one of this it tends to go into the hundreds or even 1000s of matching matrix that you have to do. And even in those cases, what you see in the lab production is very hard from lab to actually validate this is the performance it's going to have in production, and, yeah, it will drift rapidly over time in generating reality. What happened is that in this case, let's

say jackpot or any AI system they use, or actually interpreting documents and connecting that to the actual agent or chatbot or rat system, that also involves a lot of testing beyond what a tutorial out there show you, for example, out there, everything works on a tutorial like, for example, then I write, I'll just put this here. So magic happen. Let's do some chunking and answer Well, the truth is that, if you notice, a couple of those are with very simple. Enterprise documents are not that simple. And if you are an engineer, you see your documentation with a lot of diagrams that represent different parts. You need to understand which section map to another section in that book that is not naive rag for that, you need far more extraction of the context and mapping those around. So that's when you'll see that it's very nice to do a bi coded demo, but then maybe that into production, it becomes a series of experimentation which algorithm to use for this type of document versus the other Yeah, so and then everyone probably have discovered by now that each model responds quite different to the same problem, so you have to start adapting that. And there's no easy way. I mean, there's a gazillion libraries out there, but there's a new model every single week, so you're going to have to find out yet another way. And what happened in all this is that, well, there are not enough data scientists in all the teams trying to do this, so people tend to just measure what is easy and not necessarily what the customer will or the client of the application will actually be looking for. So from the enterprise data to actually making something useful, there is a gap, not only because of the knowledge, but it's a universe of combination that supposed to be done by AI expert or data scientist, but that's not the reality we need everyone to be able to do this, not on a byte coded fashion, but actually In a way that they can demonstrate this works. So this is precisely what we started, okay, tapping with it. Okay. So we know all this, and there are experts that know the right combination or kind of the right algorithm for the different type of principle or strategies. So let's start encoding, what if we start encoding that expertise on a system that can basically do a search and experimentation, let's say an experiment 100 100 experiments at the same time by just presenting these are my documents, or this is my data set. This is my use case, and this is what's important for me, and now so magic and the automation or the outdoor experience, you will go interpret that as an expert and start running all these experiments and generating the code and generating the model, in the case of traditional AI or generating, for example, the RAD system, the pipeline for your tables, the pipeline for the industry. So you then introduce those directly in your code. And when it does that, at the end, you see a leaderboard with exactly this algorithm, with these parameters, or hyper parameters, this is the performer. So you actually have a full selection of all the experiments that happen there. There are things that, for example, it's very easy to say, and I've seen this in an organization, oh, let's do this 100 tests or questions, and let's have someone go SME on this go and actually label it what should be good, or if this response is good or not, that work for a PLC. Now in data science, when you start talking about this, you have to start talking about interviewer disagreement or alignment, and that basically means that if I give exactly the same 100 questions to, let's say, three four people, there's this very unlikely they will rate it the same, because each one has different type of experience, A different type of understanding. What is the task. So that is, for example, something that people do not tend to measure because it's hard. They just have someone labeling and just use that. So for example, that's the capac score there. So this is something that an expert will know, but there's no reason for everyone else having to know that. So that's the kind of expertise we encode in this outdoor stream. Is okay? How does an expert will approach this from the data science or AI part? Yeah. So in enterprise itself, what is happening is, for example, that one of inter rater alignment, that's one that is very common because you will not see that map PLC, you will not notice that. You will notice that when you actually run 1000s of samples and actually start measuring, so what you you end up having is that without those type of evolution, there's a risk that you are introducing, and risk exposure to that orientation or to that deletion based on, okay, so this is not what it means for us on the result or things like


that. Hide the problem. Because, again, if we are just using just simple labeling without understanding, for example, the alignment across, it's very easy to say, I pass everything. I pass my test. There's not passing actual customers or cutting tests. And so the auto experiences is about that. It's about encoding all this. And have the expert there now. Also, this is a field that is changing every single day, every single week, so what you see today will be quite different from a month from now. Why? Because we will be updating these architectures and this everything behind the scenes on pretty much monthly basis. Just the idea is, again, the expertise out there encoded such that now any user can actually use this from the outdoor experiences we talked today about automl and auto red, both available today, but for example, these are some others that we are working on, identifying the right evaluation for your use case based on your industry, identifying how to measure the right risk and do, for example, red teaming for your industry, your domain, your use case, based on how the Organization defines the risk and how the industry define risk for that domain. So that's another one, which, again, is on the Aalto red teaming. And the other one is Aalto judge alignment. You probably have seen already a lot of tutorials all there that says, hey, yeah, use the model and ask the model which option is good, or is this good or bad? That works again for a PLC for but when you measure something like the Kappa score there, you will see that they tend to be given something like Opus, it tends to be about point 18 or 18% to give you an idea, you should never put something in production that is not at least point seven. So that is how different is what using a PLC versus what you need for production. So the alignment of Judge is something that is critical to actually do this rec as well. So these are some of the ones that are coming as well. And with that, we'll leave you with a DT to go over the last year you





S Speaker 2 12:39
so thanks for them for introducing auto ml and auto rack. So pardon my voice has gone after but so essentially, how many of you, true of fans, have tried building a rack system?






S Speaker 3 13:00



Awesome. Anyone here not familiar with the rack what drag is? Awesome.

S

Speaker 2 13:09

That's great. So you all know about what drag is. Drag is retrieval, augmented generation, and that is currently the enterprise AI patent for adding knowledge clear data. So like we mentioned, we are working in the connecting data to models and agents further. So like the first thing that you would do if you want to spin up a customer chatbot or customer support agent, you would do a rack system, a rack kind of pattern to be able to add your data, your domain specific data and expose that to the end user. So we are seeing increasingly that rag has become the default pattern that industries are using in generative AI. And it is used for knowledge discovery, for documentation assistance. We use it internally for OpenShift Lightspeed, if you're aware of that. So there are multiple we have a platform on our documentation called as ask Red Hat. There's actually much more than documentation. You can ask it any questions that it can tell you all the troubleshooting questions and everything available in our docs. So all of that behind the scenes is powered by a rag Chatbot. But increasingly, we see that building these rack patterns is hard finding the right accuracy, and so they've said there's a lot of like, trial and error, and a lot of like, making sure that there is a lot of testing, but people don't know where it's going to break because it is manual. So that's why we are shipping these auto patterns, and that's what we introduced for the rag. So yeah, so drag and enterprise is broken. There's no evaluation. There's just too many choices to build a successful rag application. You need to understand the right chunk size of how much chunk your data, the embedding model to be used, the chart model that you want to use. So these are a lot of variables, and if you're not an expert, as like William said, there's there's scarcity of AI expertise. We are also learning, and there are very few AI researchers in the company. So without that, it quickly becomes hard to build these systems when everybody wants to build a rack chat button. So that's the premise for these automatons. And so we are introducing auto rag, auto rag and automl. These are both available in Roy as tech preview right now, 3.4 what auto rag does for you. So it provides you this automatic, automated experience, essentially to trigger a rack optimization run. So you just provide your input data, and you need to provide some small evaluation data, and then just select the optimization metric that you want to build. So for RAC, if you're familiar, there are metrics that people care about is answer correctness, faithfulness and context correctness, that the answer is retrieved from the context that was there for the seed question. So these are some of the variables that these are some of the metrics that people optimize on. So you select the metric that you want to optimize on for bringing the language, and then behind the scenes, it does all the work for you. It uses kubeflow pipelines to trigger the whole process. And behind the scenes, it is also like using an open source library, AI for rack, that was started by IBM researchers. So now it is open source, and you can just point your Jacob pipeline to that library to be able to test out. It builds the search space and does all the hyper parameter optimization of chunk size, embedding models, and gives you the right pattern that should be deployed. So essentially, you give you give the input is just your input documents and the vector database that you want to choose and a small seed evaluation data, which we are also trying to automate for you in future. And then finally, as output, you get a whole leaderboard. So,





S

Speaker 1 17:19

so something here. So this is the first release of it. Okay, so the first release is tied to the next release. You will see that already include things like the rapid track. And the other part is, did you mention about the user having to provide the evaluation data set? Today, we have another tool, the SDG hub, ability for you to pass exactly the same document and it will generate the evaluation data set for your use case. We have not connected the two together in the automatic fashion, but we do have both the tool that allows you to create the evolution data set and the tool that will allow you to do the whole discovery of the best architecture for the use case everything is use case driven,

S


Speaker 2 18:20

and drill down on that point like that now that William mentioned SDG hub, we did a presentation on that. Also. It's synthetic data generation, so manual data, it doesn't scale right. Like you can have like 10 question answer pair, but if you are building the rack chart board on like 50,000 documents, and if you're providing 100 question answer pairs or 10 question answer pairs, you're not doing testing and you're just putting it into production. It's not going to work. So you need, like test data at scale, and that is what synthetic data generation helps with. So we have a flow for generating rag evaluation data sets, and so the key features for the current release and future releases we have a lot more support, is it does the search space preparation for you chunk size, chunking strategy, embedding method and the LLM model. So it tries different chart models and LLM models and chunk sizes, and then gives you the leaderboard here that you see of this is the pattern that you get out of it, that embedding Model X with chart model Y with chunk size a with some chunking strategy B is the best one for your use case and for your data. So that is, essentially, that's what you get out of it, the rack templates or the patterns optimization. And there's a whole leaderboard. So if you do not feel convinced, you can dig deeper. You can try other other patterns as well. And then we also show you pick results. So if you see here, this is a sample question and answer pair. So see in the input that you provided the sample question and answers, you're also able to view results. On top of that, if it is performing well on at least a 10th 20 questions that you provided. So you can do a quick scan.

S

Speaker 1 20:14

So something here, so using pattern, this kind of a link, when you click on that you get the actual code, because this generates the code only the algorithm or everything. You have the code for the ingestion, and you have the code for the regime. Because the idea is not to just assume that we can just generate everything for you, is that you can take the results here and embed those in your application as well. Okay, so that's part of what's there something we're working now is that eventually will be deployable such that automatically becomes available to the gen ed student, so you can actually do the full experience and actually run the chat application, for example, from here, launch that and automatically get the result and run side by side, right? So that's that's. Those links give you access to all that.



S

Speaker 2 21:15

Yeah. Again, this is an important point, because we are giving you a good starting point. But then if you want to still go and optimize it further, you get all the artifacts to get the notebooks you can go and again, tweak the parameters. Maybe you are, you are an expert. You have a lot of experts, and you have built 10 crappy chatbots, and now you are an expert. So, so yeah, so we provide you all all the configurations to further tweak. And this is a good this is a starting point that you don't have to like it is trying behind the scenes, maybe 1020, patterns. Another input parameter, so I forgot to mention is that you can set the number of patterns that you wanted to try. So the maximum that you can try is eight, because the more you try, the more time it will take, the minimum, I think, is cool, you at least should try so

S


Speaker 1 22:06


something real quick. For example, yesterday, I was talking with a customer that they have very engineering heavy documents and all their data more for from HR and things like that. In those cases, as was mentioned before, the optimization for the indexing is going to be different. So in that case, you can go and run and discover the best pattern for, let's say the HR documents and the other one for, let's say the engineering diagrams, etc. Now you have both indices, you have the code so you can ingest those together, but the retrieval would be pretty much the same. So even those micro optimizations now are available for

S

Speaker 2 22:55

you. So yeah, we have touched a little bit on this, but right now, auto rack is a tech preview. First goal for us is to harden it and like, get it into GA. But also we want customer feedback. So please try this. Give us feedback, and we will, we are trying to as one of the major integrations we want to have is with Gen AI playground before we GA this so that you can just quickly the pattern that you get, you can test it in our Gen AI playground, and that's another tool in the product where you can experiment different models and prompts. And then we are also adding support for graph DAG, which is something that comes up a lot in industry and with our field teams as well, when they are doing customer POCs, and we are adding mI flow, you have not already heard is something is, which is GA in three, four. So we everything that we are building observability, there is going to be ml flow. So we're going to tie all our pipelines. Everything will be tracked in mlflow, you can track your runs and mlflow, see all the data and essentially come back to it if you want to maybe do a different, different pattern in the future. So everything, everything will be available for you and multilingual is another area, because a lot of our customers across the globe, and they want to do like rag on different multilingual cortexes as well. So you can add support for multilingual as well, which





S



Speaker 4 24:30


will be harder and quality tested.

S

Speaker 2 24:36

So yeah, idea is, again, we talked about generative AI, but predictive AI, similar challenges. And this is something which this automl, it kind of automates. This is like predictive AI, similar challenges. There's a lot of models available. Not know where to start. My roommate works in finance, and she she actually was talking about starting with automl for forecasting. She works with factors and asset management and asset management firm. And it was interesting to me that a lot of like companies all for customers finance customers are still doing predictive analytics, and they use tools like automl. So we now have automl natively available in run on AI, and it is again, tech review, automl sort of does similar things for you, that it uses a kubeflow pipeline in the background and tries the search space of different models and different not another models in this course, sorry, it tries different algorithms for you, and tries tweaks the hyper parameters for you, and then eventually gives you, again, similar pattern, but this is for predictive AI, and in this case, in automl, the input for you is going to be






S

Speaker 1 26:04

so for gratification. Because if you have heard about automl, this is not the automl.org type of library. It is not that. Fortunately, we have kind of the same name this automl capability is based on auto glue on project and the experience is, again, bring your data set, define your use case, and from there we will basically, well, you have to tell us what you have to what you want to predict, and from there we go and design the experiment. So it's it's not the same as the auto not fork. This is more based on the outdoor blue one project. Just wanted to mention that because I noticed some completion during the booth earlier. So just wanted to highlight that.



S

Speaker 2 26:55

So yeah, every both of these projects are based on open source projects, autoragons and

S


Speaker 5 27:00


systems.

S

Speaker 2 27:01

Rack is based on auto. Gluon is a thriving open source project for predictive AI and so as mentioned here, you can just give your data set. You can upload your document, or give your s3 credentials for where your data is saved, and then you select the kind of prediction you want to run. So we have currently support for binary classification, multi class classification, one time series forecasting. And then it you select the target label that you want to prediction target that you would want to run the prediction against. And again, the top models to consider is like, how many optimizations you want to run, and then, essentially, you let the system handle the rest. And again, as output, you get the leaderboard where all the top models are ranked for you. So you see accuracy, you see different precision and recall source for your model. And you can further drill again. These links are again as well, as mentioned in order. It's not just like just a static video mode. It has assets attached to it. So you can go, you can drill into the metrics. You get, again, a notebook here that you can further optimize and click on. And we get tools like infusion matrix, like, if you're doing classifications to understand where your model is getting the predictions right, where it is getting confused, you also get feature summary, so, like, the relative importance of each feature that drove your model's prediction. So you get all kinds of tools to measure accuracy. We're adding more and more accuracy tools for like back testing results for time series forecasting. So all of that will be available for your data scientists, or even if subject matter experts were not technical to kind of trigger a run and get an output very easily. So I took it on, like to build, like a demo, and I got the results in like 10 minutes so, and that was like great results, 92% for auto ml as well. Our goal is to harden it go to GA but there are a lot of like things that we want to tweak. For example, in this case, you use a hard to select the algorithm and also the label that you want to run prediction on, where we want to build automations here. Also, like, given your data set, we'll identify what kind of data set it is like, this requires classification. This requires time series forecasting or regression, it will automatically suggest for you that this is the prediction point that you should run, and this is the label against it makes sense to run this. We can obviously change those things, but we have some smart, smart defaults for you. We have smart presets for you to like train faster versus train for accuracy. So those presets will also be available. Then we'll also have another thing that we are doing in this like currently, you can serve models once. If the model is built, you can serve it with a custom runtime for auto glue on. But that that has to be in three or four it is a manual process to build that custom runtime in future releases are for start. It is to make sure that is also taken care of you. You have a default runtime for auto glue on protective models by default with our with our models of capabilities in ROI and Yeah, and





S Speaker 1 30:40

so you notice here predictive model as MCB actually the demo that we have there, actually, that's what it is. Now, how much of interest is that for you guys? I also want to make sure that we're trying to overdo things here. So basically, what that does is that automatically takes the model that you have with all the prediction and wrap it behind an NCD server such that now an agent can consult that. And for example, the demo we had, there was a bank for for example, approving or pre approving for a loan, you will use operational model, but the process explaining and trying, the whole thing was more identical, and was also using a wrap for explaining the decisions. So is that a type of use case that you will find useful or not?

S Speaker 6 31:41

I feel like it's there's a lot of value there. I haven't identified a use case, but I feel like there's definitely something there on

S Speaker 7 31:48

data analytics, on infrastructure, for example, because I'm like an expert of your domains that can farm out to others that are specific to each region. I think it can be helpful.

S Speaker 5 32:08

Okay, cool. Another thing

S

Speaker 2 32:12

I want to mention here is automl. We are, it's, it's been there for some time, but we are adding more capabilities on top of that, like even Gen AI features, what William mentioned, like the ability to expose a predictive model as an MCP server for your generative AI models, to be able to call that predictive model. That is a pattern that we see a lot like it is everyone is doing hybrid AI. It is not just like just generative or just traditional. It has to be both. And when you will be building systems, you will probably want to do both, that you do predictive AI for determinism purposes, and that you do generative AI to explain what just happened. So so that is, again, on a roadmap. And then there are now foundation models for doing time series forecasting and whatnot, so we add those supports in automl as well. And then, yeah, like late 2026 we are targeting for a GA with automl, and hopefully all of us will try it, try it out, and give us good feedback. But

S

Speaker 1 33:24

want to give us good feedback, but as well, okay, yeah,

S

Speaker 8 33:31

yeah. So at this point,

S

Speaker 1 33:35

so she's gonna try to Do one of the Spirit called them about

S

Speaker 9 33:50

yes,

S


Speaker 1 34:00



oh yes.

S

Speaker 2 34:02

So like, have you guys attended the spotlight session? AI spotlight session? Okay, so some of you might. So this is the similar demo. So it's essentially, this is Fedora capital or fictitious bank. So this is federal capital, a fictitious bank that we built for the purpose of the demo where borrowers can upload documents and get loan decision in minutes. Let's Meet Sara. She logs into the portal, and she uploads her documents, and then she waits for her loan decision, and next we see James, the loan officer, logging into the portal, and he sees his part kind of applications. He pulls Sarah's application and sees that AI has extracted most of the data for Sarah, so he simply accepts, basically, he has extracted all the information he approves. But how did the system know about Sarah's credit worthiness behind the scenes? It all starts with data, and that starts with prompts. So the first way to add data to your models is through prompts. And this is the Gen AI studio that you were talking about. So you can play with our models. You can basically, in this case, we set the system prompt as a loan officer, and then you can interact with the model after your prompt unit. Then you can ask questions specifically, like, what is the current rate of fixed your mortgage? And then it is acting as a loan officer for you, because it has all the relevant context. And then, if you move next, that model that actually James was using all the data that was populated for him, that was actually built using automl, and that this is the UI for automl. So you upload your documents, you select the prediction type, you let it run, get an either vote. In this case, we have 92% accuracy. Further drill into information and get like all the tools. And then finally, we have auto rag. So a similar experience, very similar UIs, we want to provide similar experience. Upload your data, select your vector database, select the evaluation data set that we talked about, and the optimization metric, answer faithfulness, answer correctness or context correctness, and the maximum rack patterns, the number of algorithms you wanted to try, and if you can also add your models in their embedding model of the chart model. And then finally, it does the math for you, you see, you can see the pipeline steps also that's running in the background. We're actually enhancing this view for more like explainability, so that this is not a black box for you guys. And then this you get, again, a read about you drill into this pattern information. Can you go back to the last one? So this has the pattern information, answer correctness, painfulness. This is all like decent numbers, like at least a good starting point, and you can further tweak these numbers. It's very hard to get even 91% accurate is answer directness value with rack systems. So it makes it really easy for people to just deploy a rag and get started with rack systems. This is you get all the other information, chunking and letting retrieval, generation, and then finally, the sample QnA, the data that you had provided, the evaluation data set, you can see that What documents do I need if I am self employed? And it's able to tell you we need the last two years complete tax returns and all the other documents if you're self employed. So so this helps the loan officer also, and also helps the users like Sara to understand how to give them documents. So that is, that is the experience we have with automl. And you can actually look at this, well, we also have a question, question, the





S

Speaker 1 38:19

demo, just run through that URL. You get the click through experience, but the actual demo, it is as an AI team. Start. Okay, quick start. Okay, so you can actually go and get the code for that. So that's on this side. On the other side is we really want to so for us, it's not just about the tools, because as a product manager, my philosophy is the best technology is the one that the user do not see. It's about the experience. So we really want to understand what kind of experience you want to see there as a user, as a field essay. So if you are up to it and you want to participate in the US research experience, please also sign up for that, because we really want to provide the best experience possible. For me, it's not about the tools. I take care of that with my teams, of making sure the signs is sound, but I want the best experience for the user.

S

Speaker 2 39:27

Yeah. Also, if you have any questions, let's I guess

S

Speaker 6 39:33

I had a few questions. First one, well, I noticed that there's a safe notebook option. So at the end of the auto rag. In the auto ml, you end up with a notebook that you can modify. Okay, of course. The other thing is, in the auto rack stuff, who or how are the patterns being defined, and just high level, like the chunking strategy, the embedding models, and then, and then the models that are being used to, like, actually answer the rag questions and stuff like

S

Speaker 1 40:06

that. So all those are parameters that you can select. And for example, you say, choose and try 345, embedded models. And it will try to find the best embedded model within the universe that you provide models in the future, we want to make that even more dynamic, especially for very domain specific areas. We want to be able to automatically fine tune the embedding model for that domain such that your similarities errors and vector search becomes far more accurate.

S

Speaker 6 40:45

And how, you know, I haven't really got my hands on Red Hat AI yet. I'm hoping sooner in the future. But how well does the patterns that you end up defining? Let's say you're happy with what you got from auto rag and stuff like that. And then you want to take that data that's coming from that rag solution and tie it into an agentic workflow and all this other stuff. And then from there, how well does that? How well, how easily is it to tile this back into things like trusty AI and ML flow?

S

Speaker 5 41:20

I so

S

Speaker 1 41:21

that's very easy. First mlflow, native mlflow capability for experiment tracking that is actually coming for three, five, just we did not have enough time to complete that now, but that will be there natively to answer the question, and also your previous question about the templates. So this will have a template mechanism that you can even define your own architectures. So as part of all the algorithm and the architecture for rank, you'll be able to onboard your own as well, such that they are account as part of the experiment, because that's how we want to make this very pluggable and easy for any expert or anything out there to go and bring a new option, and we will consider it as part of the use case when it match the use case. So the other part, on the safety side, all that is bringing to it throughout the platform right now, something that I don't know if it was here, is that that will integrate with eval hub. And eval hub is our AI evaluation control plan, and from there, you can run any type of evaluation, even AI risk or traditional benchmarks or guardrails, even. And from there, we can also go and evaluate, for example, the guardrails for this. And if you want to connect guardrails to this, we are creating, actually a standalone guardrail system as well, such that you can plug it at any place. Today, you will have to use it either through the AI gateway or enabling it through trusted

S

Speaker 5 43:11

AI

S

Speaker 1 43:11

just in case. The only reason I asked

S

Speaker 6 43:15

about those two specifically is it feels like everything I've learned this weekend or this week by default, I probably wouldn't want if I want something in production, ultimately I'm going to want ml flow and then ultimately guardrails as a service, no matter what pattern I end up deploying.

S

Speaker 1 43:32

Yes, so something there, and that's more on the evaluation side, is that once you have something in production and you enable tracing from the traces, we will be able to do continuous evaluation and then flag you when things happen,

S

Speaker 2 43:48

and even it's not like that, this isn't used use for the current state, because cube for pipelines also provide you some it tracks things for you. So experiment tracking by nuclear pipelines also provide, but natively we want we will have mlflow By next few weeks. Okay, cool.

S

Speaker 1 44:08

Thank you very much.

S

Speaker 5 44:09

Thank

S

Speaker 10 44:09

you very much. Yes, I hope you guys have

S

Speaker 1 44:12

a wonderful day. Bye.