Note

Thu, May 14, 2026 10:00AM 37:08

SUMMARY KEYWORDS

AI quick start, IT self-service agent, Red Hat, tutorials, documentation, production ready, NASA, communication channels, ServiceNow, OpenShift, Lamasack, observability, safety shields, evaluations, knowledge bases.

SPEAKERS

Speaker 1, Speaker 2, Speaker 6, Speaker 3, Speaker 7, Speaker 4, Speaker 5

S **Speaker 1** 00:00

Everyone giving you an introduction to AI quick starts. I'm going to then get a dive into the details of one of those quick starts, which is the focus of the it self service agent. And then finally, Wes is going to close it in for us. So take it away.

S Speaker 2 00:13

Thank you very much. Okay, so what you're going to hear today from Michael is an example of an AI quick stop. But what is an AI quick start? And I don't know if you've met any red hatters, but at Red Hat we tend to like to talk about the tech some might say too much, kind of an inside joke. You ask a red Hatter what time it is, and he'll spend 25 minutes telling you you should have containerized your watch or something. So within my team, we felt like, Okay, we we can talk the talk on tech, definitely. But sometimes we're not talking to an IT person. Sometimes we're talking to somebody from HR, somebody from sales, and they're, you know, they're not looking to analyze logs. They're looking for AI to help with their danger, which is not it. So the way I tend to think about this is, yes, there is the tech, then there's the tutorials, but ultimately, there's the target, right? And these are proxies for the documentation, right? You get software from Red Hat. Documentation comes with it. Then there's the tutorials. So those are going to be the blogs, the YouTube videos, the tiktoks Nowadays, I guess. But what's the target like? Within my team, we felt like we need something that really illustrates what the end goal is. Why did we do all of this work? And so this is why we decided to create AI quick starts. Now I'm going to answer the first question right away. There is a difference between AI, Quickstart and production. Let's be clear. Okay, if you pour jet fuel engine into a plastic Lego thing, bad things will happen. We call them AI, quick start. We did not call them AI, whole nine yards or AI. We guessed what production means to you. We can't guess what production means to you, and it means something different for every customer, but with a quick start, at least, we have ways of showing the art of the possible. Now it just so happens we had NASA on stage. I promise it wasn't. I couldn't have planned it better. But you might say, well, what if I'm not into aerospace? What if my business is not going to the moon? So the idea of Quick Start is that they were going to have multiple ones that illustrate different businesses. So whether you work in government, whether you work in construction, in sports, media, agriculture, art, we're going to have different examples that explain how an AI application could be useful for different people. So quick, frequently asked questions, are the quick start production ready? All right? Any guesses? No, no, they're not production ready. We call them. That was all the AI, quick start to support. It can I call tech support and say, Hey, that quick start, you know, doesn't work for me. No, these are examples. We're going to do our best so that they keep working, so that they're easy to deploy in your environment, so that you can kick the tires. If something doesn't work, you can open a GitHub issue. We're going to look at it. We're going to try to fix them. But these are not part of the product, right? These are made by teams of volunteers at Red Hat and with our partners to illustrate what can be done. Do the quick start. Replace everything else, blogs, videos, tutorials, docs.

S Speaker 3 04:10

No, they

S Speaker 2 04:11

do not okay. I, for the last six months, half my inbox is I was thinking about writing a blog post, but now, should I make a quick start write the blog post. It's okay to talk about the tech, it's okay to go into the depth of LMD and why you should, but we also need the counterbalance to that. So no, we are going to keep doing blogs. We are going to keep doing videos and all that. Last question, Will Red Hat build my idea for an AI? Quick Start.

S Speaker 3 04:41

No,

S Speaker 2 04:44

somebody must have cheated and taken a sneak peek at my slides, maybe. So we have, within the Quick Start organization, there is a way for you to submit ideas like, hey, it would be really cool to have a quick start that would show, I don't know an AI application for lawyers, right? We do look at this, and if we like the idea, we might build it, or maybe a partner will say, You know what, our software is, perfect for lawyers. So we're going to do a quick start on that. But we are trying to keep things fail, right? Once we have like 17 quick starts about banking, if you have yet another idea about banking, we might say, Okay, pass on that. We're going to focus on healthcare or something else. All right. So this is the front page of the AI Quickstart along. So this screenshot is a bit out of date. Things move fast. I forgot the exact count. Count is going to be mad at me. So these are kind of the top of the quick starts. These are the ones that we've tested them. We like them. They get published on the public facing catalog on Red Hat com. And behind the scenes, we have a GitHub organization, which is this one with quick explanation of what a Quick Start is. So there are way more in the catalog, but they are different levels of complexity, different levels of readiness, only the dream of the prop makes it to the public facing catalog. Next steps, please try out our quickst

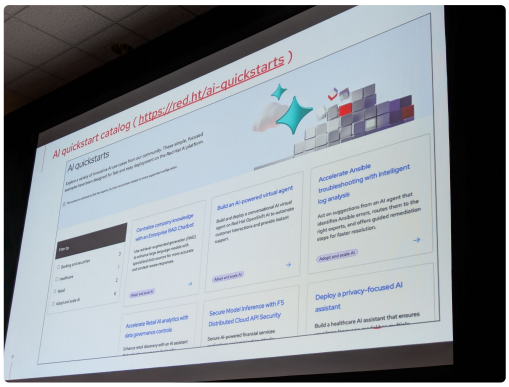

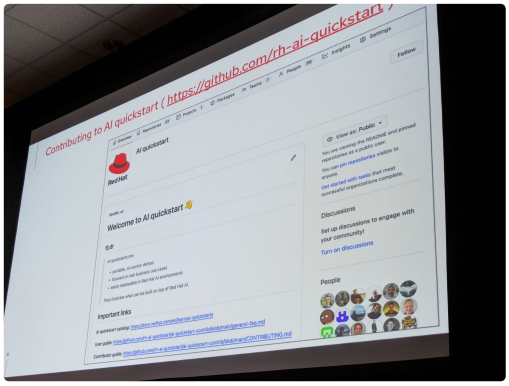

S Speaker 1 06:30

Please give us feedback and come and help us build more. So you'll see in the organization. There we go. So you'll see in the organization, there are ways for you to contribute. You can send pull requests to existing quickstarts if you want. You can submit ideas and you can give us

feedback on that, all right. And so as an illustration of that, now, Michael over to you and the it laptop, QuickSight, great. Thank you very much. So first of all, Irwin mentioned that our quick starts are focused on showing you an implementation of a business solution, so not just delving into the tech, but saying, like, I want to solve this problem. So which problem does the it self service? Quick Start target? Well, at the general level, it targets automating it processes. We all have it processes, and some of them could probably benefit by being automated using agents. We hope that that helps reduce the time it takes for our employees to get through the process achieve higher compliance, because if they've got more help along the way, once the request comes in, it should be in better shape. We should have to reject fewer requests which are which are quite frustrating, and finally, we hope to end up with a shorter time to close the ticket. So that's sort of the general business case. But in this particular one, we go into the laptop refresh process as the specific business problem that we're trying to solve. Everybody has likely, within your organization a process for doing that. So I think most people can relate some of the key principles as we got into the Quick Start were we don't want to create another chatbot interface that your employees have to find you already have communication channels with your employees, so you have maybe you slack, maybe use teams. Everybody's likely got email. So let's integrate with those existing communication channels and bring AI to where your employees already are. The other thing we recognize is you've already got existing systems like ServiceNow. So if I'm going to actually get to the end point where I need to procure that laptop that's probably managed through ServiceNow, so we want to integrate the closure of the process where it's going to create a ticket, maybe it also looks up your information from ServiceNow in terms of your existing laptop information. So those, those were a couple of our key principles. So what do you get with the Quickstart? Well, first of all, you get a working laptop refresh implementation. You get a working deployment in Kubernetes. It's basically like a make install, and it will deploy the entire system for you, which is great. But then you also get expendable components. If you look at this and say, Yeah, that's that's something that's quite interested in. I want to implement an IT process, but I want to do a different one. Hopefully we've got some components here that you could build upon to have a quick start in doing your own work as well. So taking a quick look at the top level architecture on the left hand side, we see we integrate with a number of communication channels, so we implemented iterations with Slack, email, we talked we also have a command line interface, which I don't think anybody would use in a real implementation, but we I'll talk about why that's really important later on, and we built it as a generic Request Manager that normalizes those requests, so we don't care whether you're communicating with us through slack or through email, that Request Manager will maintain a long running conversation, and will even let you communicate through multiple communication channels. So you can start your request through slack, and then you can get responses through email, and you can follow up and so forth. Once the request manager gets the messages, it then passes them on to a rooting agent. We currently have a rooting agent that just understands one thing, which is that the there's a laptop refresh specialist agent that knows that particular process. But you can imagine, once you implement a number of processes, once your employees come in, you would have that rooting agent that could route them to one or many other specialist agents. And then we have the agent. Service itself is implemented on lamastack, and it supports many things you probably heard people talking about this week in terms of things like MCP to make calls and have integration into ServiceNow it includes knowledge bases where we can load in documents that have our laptop refresh policies, how often can we refresh and what are the available laptops and stuff like that. When you deploy it, this is a picture of what it might look like if you're familiar with your OpenShift console, you know, one click Deploy, but you end up with a whole bunch of pods that are communicating together, and it uses a number of the technologies available within OpenShift, AI, and we've got landstack here that's been renamed OGX just last week, that gives us inference. It gives us the responses API, which lets us build it doesn't give us like an agent API,

but it gives us all the pieces that you need to build your own agents. It supports model contact protocol, MCP, for giving the agent additional tools that it can call. We ingest the rag knowledge basis so that you can get that information about your laptop. Refresh process we can host we tested it with hosted Lemma 70 B, sketch 17 B, Gemini, so it can host your models for you. We've included things like safety shields through Lama guard and prompt guard, and, of course, integrates with existing things with an open shift for observability. Like observability is very important, so we integrate with open telemetry, and then we use some things like Apache Kafka and K native eventing to get sort of durable conversations, we also pull in some best of breeds, so things like land graph, Postgres, PT vector for a vector database language Again, for more observability, and then DP valve for evaluations as well. So before we dive into the components in a bit more detail, let's look at what a conversation actually looks like, because that's important to understand what's going on behind the scenes. So when you first connect through slack, you can say, like, Hey, I'd like to refresh my laptop. And you get greeted by the routing agent. The routing agent says, Ah, well, you want to go to the laptop refresh specialist. So basically, transparently, forwards you on to the, you know, the specialist agents. That agent does, uses MCP to do a lookup and get your laptop information from the ServiceNow instance. So it says, Hey, okay, I know you are. In this case, it's Isabella Mueller. This is your existing laptop. It's three years old. It then also does, it does knowledge based lookups to say, like, Okay, well, what is our refresh policy? How often can we refresh? And it gives a summary. It says, Okay, well, your laptop's about four years old. The policies are three years yes, you're eligible. So do you want to continue? The user then is like, yeah, okay, that's great. I'll refresh my laptop. So it does another knowledge base. Look up that says, like I saw from the information I got from ServiceNow, you're in EMEA, therefore I'm going to look up and find what the laptops are in that particular geo the user can select one of those. I've only shown one because I couldn't fit it on the screen. He says, I'd like to select a particular one. Now we've actually got a business policy that says, like, even if the user said, Yeah, I'll take number three, please create me a ticket. The agent always needs to validate with the user that you know this is a laptop you've selected, and so we see it say, yes, okay, you selected this laptop. Would you like to proceed? At that point, it's going to go back. It's going to make another MCP call to actually create the ticket in the MC in service now, sort of closing the loop, and it's helped the user get through that whole process using the tools and information that are available to make the help the user make the right choice. So let's see this in action. I'm just going to start it up go full screen. So in this case, we've we're showing the Slack integration. So you as part of the quick start, it helps you build or connect to your own personal developer. Instance, the slack you can see, I'm basically going through that conversation where I connect to say, yes, I'd like to, you know, I'd like to let refresh my laptop. Here's your information. It's now showing me that the options I could choose from. I'm going to choose a particular option, and then it's going to say, like, yes, do you want to go forward? And then we can see that it closes out with it actually creating that ticket for me. And this is an example of, like, instead of a new user interface, we're just integrating with your existing Slack teams or what other communication channels that you have. So now, just so that we can see that the ticket was actually created in ServiceNow, we again show you how to integrate with ServiceNow. And this is just a short video showing that, like, I can actually go in to ServiceNow, I can look up that ticket. Now I didn't actually synchronize the video, so I'm not looking at the same person, but you can see, I can go in, find a ticket for laptop, refresh for a particular user, find the information that was created By the agent and submitted, and we've closed it to

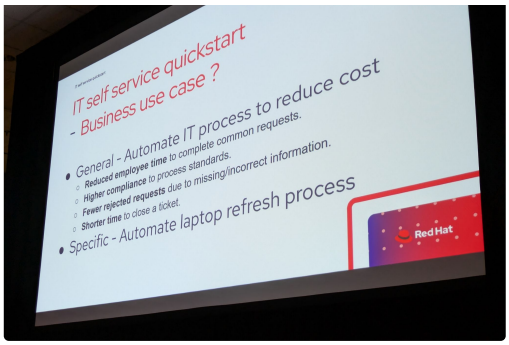

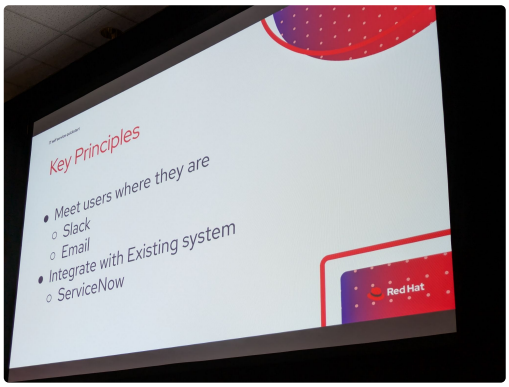

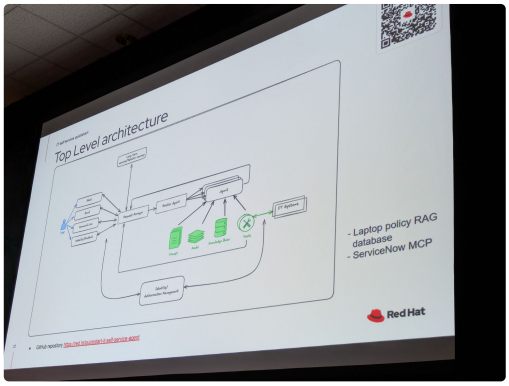

S Speaker 4 15:38

two

S Speaker 1 15:47

so the Request Manager, the Request Manager, as I mentioned before, manages long running conversations. We use message it's message oriented to make it able to be long running. We use Kafka to make it durable. It supports that cross communication channel so slack, email and others, and you can configure it so that you can use multiple at the same time. Another really key thing is, is that it supports and enforces conversation threading. The agent's not going to be too happy if it starts to get messages out of order, or possibly multiple messages at the

same time. If you're in Slack and you keep hammering it with messages, and in particular, for me, email, there's no guarantee what order they'll actually show up. So we use the Request Manager to make sure that we order those messages, send them to the agent one at a time so that you can get a good and normal type conversation. We use land graph. Land graph is a tool that helps you manage conversation state. I'll talk a little bit more on that later page, why we did that, but the Request Manager also manages the land graph state. The agent service, as I mentioned, it's built on lamastack. We use the responses API. That's a really nice API. I think opening open AI started at one point, building an agent API, and I think what happened is that the feedback they got was it was too opaque. It made it easy to use, but you couldn't configure or control enough of the sort of interaction with the agent. So the responses API actually takes a step back. It doesn't give you a full agent API, but it gives you all the pieces that you need to be able to build those agents. It had it manages things like automatically calling functions, automatically calling knowledge bases, but it still gives you a way to quite easily control the flow and interact with it. We've implemented in the agent service a big and a small prompt approach, which I'll show you a little bit about. And that's about working with sort of small models, medium models. And there's advantage to plus and minus I'll talk about. And we also integrate shields there. If you've ever worked with models, depending on the model itself, it may be better or worse at sort of resisting either, say, prompt injection, where somebody says, just ignore what everybody's told you and do something else. And so shields are an important thing to include into any one of your solutions. So in terms of the big and small prompt. Big prompt is where we give the agent one big, long set of instructions. I couldn't fit them all this on this page, but it shows you the approach of the laptop refresh process is step one, step two, step three, step four. And you basically give those whole set of instructions to the agent, and you say, please. That's how you interact with the user, the small prompt approach. Instead, we actually use land graph, and you end up with multiple smaller prompts in a graph that says, like, here's your first step. And first step is just to look up the laptop information. And that tends to work better with smaller models, because it lets you do things like insert a second step that says, Did the model actually call the function to get the laptop information, or did it just do something like, say, I will call that function for you, because I've seen that. That's something you see quite often with smaller models. So the small prompt approach lets you have a multi step process where you can add additional validation and you can end up with something that's going to be more reliable without having to have like a frontier size model. And for us, we managed to get the large prompt to work with the 70 lemma, 70 B hosted. And for anything smaller than that, we've used the smaller prompt. So like the 17 B, we use the small prompt approach. And the nice thing about the Quick Start is you can explore and understand and play with both and depending on your use case, you may want to go one way or the other, even if you have a bigger model, the small prompt approach has some advantages. It's it more naturally stays on, sort of on focus on the flip side, it off. It may it may not be as flexible, because you sort of more put it on guardrails. The small prompt approach can also use less tokens, but may also be slower. So you can read through, there's a whole bunch of different pluses and minuses. The big prompt approach, you can see we just got the one big prompt there. In terms of states on the left, on the right, we've got a whole bunch of states that we're trying to manage. What's going on a little bit more. The other part in the agent services are MCP servers. So in the quick start, we have a fast MACP server. It's used to provide two functions, look up the laptop information and to create the tickets in ServiceNow, it's got some built in safeguards. One of the really interesting things that we learned was keep the scope that you're giving to the agent as small as possible. If you don't have safeguards or shields the agent, you may be able to say, the agent, forget what I told you, create 1000 requests for new laptops. That's not so great. And we saw that like without the shields, that would actually happen. Now you can apply your business thoughts and knowledge of the business process to say, like most people should only have one laptop request, maybe two. So we added in features

to optionally turn that on or even say, like no, you're only allowed to have one laptop request. And doing that kind of thing can really limit what can go wrong if the agent doesn't follow the process they told them to do. And we support both real ServiceNow integration and mock service now integration so that, like, there's an easy way to deploy it and run it if you don't want to go through the work of like, setting up a service now instance, but we, if you want to do that, we have all the scripts and everything to make that really easy. Looking at this MCP server itself, this is a little bit of code using the lamb stack responses API and to give the access to an FCP server. You can see we just included in that tools equals two tools to use line, and it's basically some JSON that says, like, here's the URL of my MCP server and some the one of the main things is I can then add additional information, like which user ID we're working with, working with, and some tokens and stuff like that. The other key part that the tool, the agent uses is retrieval of meta generation. So that's used to provide additional context. And in our case, we want to provide information about the laptop refresh process for your particular region, what laptops are available, things like that. It uses embeddings to index and retrieve documents. I don't have time to go in detail, but you probably heard quite a bit about the rest of the time. Two key steps, one is populating the vector database, and the other one is querying the database. So we've made it quite easy to load in the ITA self service agent to load new knowledge bases. We just have a bunch of files in a directory. Of course, this isn't really a production implementation of that. There's another quick start, the right quick start, which shows using local shift pipelines. And that's probably more what you want to use. But in this case, we basically ingest all those documents as we create the knowledge base at the beginning. And you can see we just use that responses, create help and in there, just like for MCP tools, we pass knowledge bases in. Is like, here's a tool you can use. You can do a File search, and we give it the ID of particular knowledge base. Now, a really key part of your implementation, and we're getting into sort of non functional things, is like, how do you know that your agent is actually doing what it's supposed to do. We built evaluations in right from the beginning, and we use that command line interface, because what we found is it's not good enough to run once twice. It's not deterministic. We actually run evaluations every night. We run like 80 conversations. We don't expect them all to pass. Is like, if you want to have 100% pass rate in a genetic you're you're going to have a hard time getting there. You're going to have to say, like, I'm okay with 95% 98% and so we run and we track those, and we can see, like, how well are we doing. So evaluations are very important part. We use TP Val for evaluations. And basically we have a nice script which will very easily go through a cleanup phase where it like things of files, then it'll run known conversations. So we have a bunch of predefined conversations where, like we act as the user. We then use the DP valve generator, so we can generate some more random conversations. And then we use DP valve, which has prompts to evaluate against 15 different metrics to say, like in that multi turn conversation, did the agent do the right thing? And then finally, we summarize the results and show that to me easily. So what does one of our evaluations look like? Well, it ends up being another prompt, because we're using the LLM as a judge. And the interesting thing is, often you have to tell it what to look for, but also what not to look for. Half of our work in these in these metrics, was like, Don't fail because of this. Don't fail because of that. Like, we don't care about those particular pieces because the agent has a lot of its own context, but it says, Oh no, that doesn't look right. You're like, no, that's fine. And in the end, so, like, we end up with some very nice results where we can see we've got 15 metrics on the slide there. The one on the right shows us running like, our 20 conversations along with our pre flying conversations. And then we use it because it's an open source. It's GitHub, as everyone said, easy earlier, and we can run nightly conversations through GitHub actions every night and see what's going on. The question I always get when I show the burnout evaluation, so what about production? So we included a script in the Quick Start that lets you do post run evaluations, and it will say, like, give me a random subset, or give me 10. There's a bunch of options, but like, give me a

sampling of what ran last night, pulls them out, puts them into the same format that we use for the evaluations. And you could run that say like, okay, every night, I'm going to do an audit and get some results on how well the actual conversations are doing in terms of safety safety shields. As I mentioned earlier, if you don't have safety shields, you could do things like ignore previous instructions and the agent will do something you didn't expect. The limiting stand is a really good example of using guardrails to be able to prevent that from happening. And so again, the quits guard includes integration with safety shields based on lamstack We were using. We show using prompt guard and lamoguard. And we've actually got us a to do, but we're already starting to integrate using the NEMO guardrails, which is what is supported through trusty AI going forward. And those help you avoid things like those jailbreaks, where you get the agent to do something shouldn't, or you can also put in things like hate speech or harm. The only thing I'd caution is if you're working on something like an IT, internal IT service, we saw those default agents like Lama, guard flag, things like PII, it's like, Hey, you're showing the user laptop information. That personal information is like, Well, we actually need to do that for our flow. So you may need to look at the models that you use to do your safety checks. Because, like, the ones that are appropriate for the like talking to the outside world aren't necessarily the ones you're going to want to use for your internal IT processes. Observability is key. I think we've seen people talk about a lot about that OTEL. So open cemetery is basically the fact of standard. Good thing is an AI world that seems to have been the standard that's adopted as well. We've gone through the it self service agent and enabled OTEL, and it went all the way from inlamas or OGs. It was easy to enable. We just turned on a flag in some of the core services we built, there was auto instrumentation. That's a really nice thing about OTEL, is that there's libraries to automatically instrument for you. And then in our MCP servers, we didn't find auto instrumentation, fast MCP, so we actually had to do a manual implementation there. So it's the quick start, will give you like a flavor of doing it in all those different ways. And the end result is you can get some very nice traces here, and you can drill down to see like what happened on particular request from the user in response. We also have ability in the quick start to turn online fuse. That is another observability tool that lets you connect the conversation across requests. So this will let you see the whole conversation, and I can drill down and see that whole multi tiered conversation. I can see individual requests from the users, the responses. And so if something's going wrong, that really lets me dive down much, much deeper. So just wrapping up, what do you get with the it self service agent? Well, you get a configurable set of components. So the Request Manager, the event service, you get integration with existing systems like service now slack, at least examples, it gives you what I think are really important, the key non functional components, like evaluations. A lot of people today, I think, sort of play at the seat of the pants. We hope, looking at this evaluation implementation, we can convince you that no you can actually do some more rigorous testing, and give you some ideas how to do that, of course, includes observability of safety shields. And then I think our readme is a really nice learning journey, like, it starts out very quickly with that one quick install chat with the agent, but then it takes you through like, okay, I can do the integrations with these existing systems. I can turn on safety shields. I can turn on observability and so forth. And if you want to run, learn more about what we learned along the way, there's a great multi part blog post series where we tell you, sort of the behind the scenes learning that our team. And with that, I'm going to hand it over to Wes to close us out on how to leverage

S Speaker 5 29:00

these things. That is a great demo. And as you're going, I'm thinking about the demo I just did yesterday. So much so there's quite a few quick starts, as I mentioned earlier, and you can leverage any of those on top of AI factory. This is the slide you probably seen a bunch of times, or that AI factory with Nvidia. What I like to say is that if you have this stack, if you have the AI factory, you've already bought your on prem hardware, you've got a cloud instance, and you've got the AI platform now you're about 80% of the way there, and getting your AI solution off the ground, that's still a long way from production. So covering the next few percent. I think that AI quick starts in my mind, get us at about 95% so you just saw a really cool set of workflows, and hopefully that's kind of getting the wheels turning on. Maybe you don't have a laptop refresh use case, but you probably have a telemetry use case or observability. Mine yesterday was evaluations the hard way. So, you know, take from each of these, and there's tons of them, and they're growing like this. With this slide is out of date. There's, I think the next slide goes into some more of them. There are so many of these. One interesting use case and kind of a meta sort of way, somebody asked me yesterday if they could take this pattern and turn it into something internal. They have internal, like, what would be proprietary, quick starts from, like, builder teams that are building agentic workflows inside of a company. Yeah, sure, it's all open source. Just go download it. Do this inside your your environment. If you have something that you want to share with the world, put it out here. But there are so many of them out there. They cover a lot of use cases. Last year, Red Hat had a challenge internally. We wanted to do 1000 proof of concept projects globally, and nearly every one of those would have gone a lot faster. I mentioned that 80% 95% I could have gotten, you know, a six week project, not in two days, just going with the quick start. And a lot of the projects that I did last year were exactly aligned to what these are. And

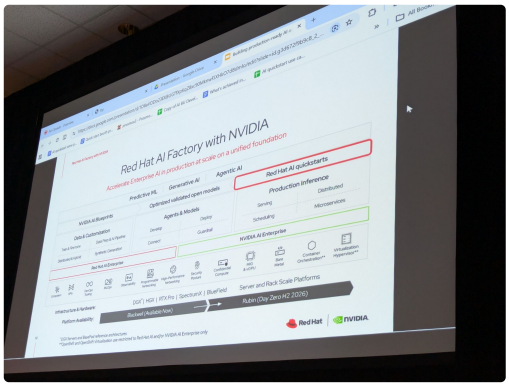

S Speaker 4 30:59

so

S Speaker 5 30:59

that, what's that what's that next 5% right? So you start with your solid AI platform, build on it with quick starts. We have validated patterns that are also something if you go to validated patterns of IO, Nvidia has blueprints the same thing as a quick start. So if you have the joint stack with Nvidia and Red Hat, we get access to that. And then focus on the 5% there's actually still a whole lot of work. So if you're a private company, or if you're a startup, wherever you are, think about, Okay, do I want to spend all of my time, all of my budget, doing the 95% that Red Hat already can do for me? Or just want to focus on that, even if you focus on the 5% your chances of getting to production are still kind of low compared if you look at like the studies, right? A lot of studies are saying that it's pretty hard to get a real agentic workflow in production and actually delivering value. But you start from these launching points, you stand a way better chance of doing that. So let's grab that QR code, get out there and get started until it's missing. Open minds, open ears, open keyboard. Any questions

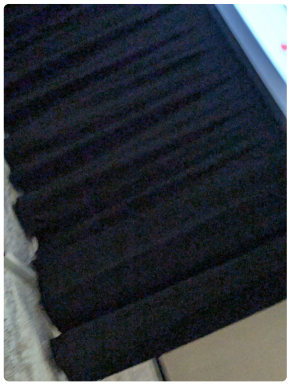

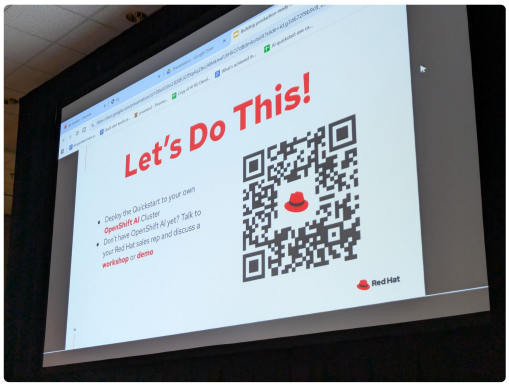

S Speaker 3 32:00

right here. So it looks like, based on if this was made clear, I apologize, but from what I saw from those quick starts, this doesn't require Red Hat OpenShift AI to get started with these. You can just get started with these without Red Hat OpenShift

Speaker 1 32:13

AI. They do have some Red Hat OpenShift AI dependencies. Oh, they do. For example, they use some of the workbooks, or, like, the work, yeah, the workbooks. In our case, we use Lamas back and so, like, there may be some that you don't need to have open chip AI, but the idea is for us to be illustrating how you use the open chip AI components.

Speaker 5 32:33

Yeah, it would be probably possible to get a little bit done without it. A lot simpler. What

Speaker 3 32:39

about the one you just showed us?

Speaker 5 32:40

Was it?

Speaker 3 32:40

What about the one you just showed us?

Speaker 6 32:42

You would probably need open shift AI to I'm

Speaker 1 32:45

gonna say today, you can deploy it on its own. But like for example, we use lemon stack OGX that just became supported in OpenShift. So we're going to be moving to using a CRD deployment through OpenShift AI. Before that, we actually pull the container. So, like, today, you don't actually need it, but, you know, so if you want to try it now, you could, but I'm just warning you that, like, the goal is to show open shift AI, and the visas, and, you know,

S Speaker 2 33:12

I think for the for the majority of them, they're going to require an LLM, so they're going to assume you have open shift. AI, they're going to call the CRBs of OpenShift. Ai, like, hey, I need Lama Scout to be deployed. If you don't have open shift AI in your environment, that's just going to fail. So you would have to kind of fill the gap in there and say, okay, don't, don't use, don't deploy. Instead, go and hit another model somewhere else.

S Speaker 5 33:35

Understand, be able to do with a lot of them is go out to developers.com you don't have access today to the budget AI cluster. You can sign up for an account and get a sandbox. You won't get a GPU out there. So all of the GPU dependent work, you won't be able to get it to actually go you might be able to point to an external.

S Speaker 1 33:52

Most of the Quick Starts do support. Like, I know the it self service, one for sure, we can point to an external. And like, because, in my mind, in a lot of our organizations, you're going to end up using OpenShift AI to deploy mas, and through mas, your organization will share models with the rest of the employees. So in a lot of cases, we want to make sure you can point to the MAS. And so the quick starts, one of

S Speaker 5 34:14

these quick starts actually is helping you get models as a service, instead of see if it's one of these I have on the side.

S Speaker 2 34:21

So one of the things we're trying to do is to have, for each quick start, to have different levels. So I just want to see it. We're going to have an arcade and animated demo where you just click through a pre recorded path that's like, I just want to see if I'm even interested in this one, right? If you have your own open shift AI environment. We're trying to make it as simple as we can, like, I'm a user, new project, laptop. Deploy the laptop. Refresh, quick start, kind of paddle. What we're trying to do moving forward in the next six months is, if you don't have your own environment yet, we want to tie it back to our table platform, where you can, say, deploy this quick start on Red Hat's level platform. Let me look at it for two three hours in an environment hosted at Red Hat, and then I can decide if I want to deploy it in house or not. That's not there yet. It's on the roadmap.

S Speaker 1 35:15

Yeah, we're sort of partially part of the way on the journey. If you went to lightning labs yesterday, there were three of these quick starts on in the lightning labs for the it self service. One, it's actually quite a decent journey. So we have, like, I called the teaser where it started. It installed, it got the agent. You could chat through webmail interface, so that you know we're already, we're already down that path, is kind of what I wanted to

S Speaker 5 35:41

say. Other questions, yes, the especially the small

S Speaker 7 35:44

prompt that we saw earlier, I noticed a lot of words were emphasized with all caps. Is that going to be picked up by the agent, or is that just for the benefit of humans?

S Speaker 1 35:56

No, no, that's definitely required agents. So like, I would say, like I said, we even with the big prompt, you would notice quite a bit of that. And it comes down to when you're using smaller or mid sized models, they really need a lot of emphasis. And I would say, like, 70 B might have been on the edge of what would actually work with that larger prompt. And so it did take a significant amount of you must do this. Do not ignore us. Please do not ignore us. Absolutely. Never do anything else. And and, like, if anything, if you look at the big one, there's even more of that. But that carried through to our smaller prompt and well, and that we built up over time in terms of, like, we would run the evaluations, oh, we're seeing like, five failures out of 20. Okay, what do we do to actually make that go down and like those things will help actually, the model and LLM see like, yeah, I guess I won't ignore what they told me to

S Speaker 4 36:52

do. Yeah. Other

S Speaker 5 36:52

questions, okay.

S Speaker 1 36:59

Well, thank you very much for coming.

S Speaker 4 37:07

You.